

Invoke Dynamic

By Mohit Kumar

JRuby Implementation

- Highly Dynamic
 - Dynamic Typing and dispatch
 - Lookup by name only (Single)
 - Mutable Classes
 - Blank States filled at loadtime
 - Closure and cross frame data(\$~, \$2, etc)
 - eval

Doesnt the JVM do this internally?

- Internally yes.
 - Majority of calls are dynamically bound.
 - Hotspot makes them static fast
- Externally, Not before invokedynamic.

Caching Callsite

- Monomorphic inline cache
 - Cache method +class serial number
 - Compare serial number to guard
- Eliminates costly hash lookups
 - No JVM can inline through it.

What JVM sees


- Ruby Code calls Caching callsite
 - Monomorphic
- Caching Callsite looks up calls method
 - Megamorphic
- Lamda will have the same problem
 - And maybe the same solution

invokedynamic

- Allows custom code for call protocol
- JVM inlines through it
- The JVMs new “Golden Hammer”
- Jruby Masters Today, 1.7ish soon
 - 2-10x Performance Improvement
 - It really does work



invokedynamic





bootstrap method

invokedynamic

bootstrap method

method






invokedynamic

ta

bootstrap method

meth



invokedynamic

ta

bootstrap method

meth

invokedynamic

ta



The deal with method calls (in one slide)

- > Calling a method is cheap (VMs can even inline!)
- > Selecting the right target method can be expensive
 - > Static languages do most of their method selection at compile time
 - > Single-dispatch on receiver type is left for runtime
 - > Dynamic languages do almost none at compile-time
 - > But it would be nice to not have to re-do method selection for *every single invocation*
- > Each language has its own ideas about linkage
 - > The VM enforces static rules of naming and linkage
 - > Language runtimes want to decide (and re-decide) linkage

Example: Fully static invocation

> For this source code

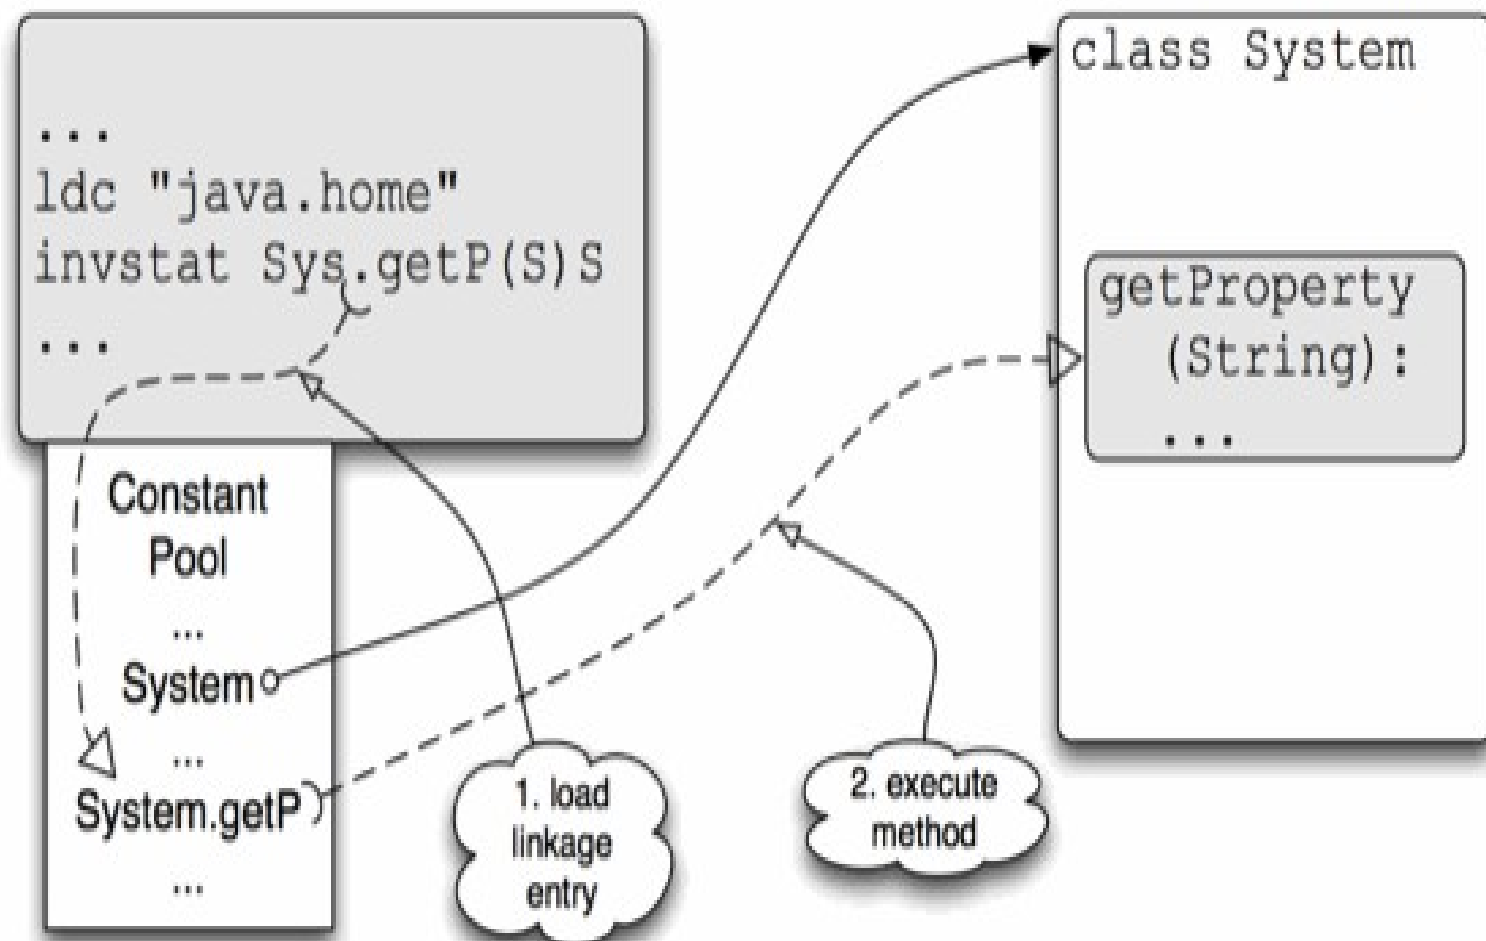
```
String s = System.getProperty("java.home");
```

The compiled byte code looks like

```
0:   ldc #2           //String "java.home"  
2:   invokestatic #3   //Method java/lang/System.getProperty:  
                          (Ljava/lang/String;)Ljava/lang/String;  
5:   astore_1
```

- a) Names are embedded in the bytecode
- b) Linking handled by the JVM with fixed Java rules
- c) Target method selection is not dynamic at all
- d) No adaptation: Signatures must match exactly

How the VM sees it:



(Note: This implementation is typical; VMs vary.)

Example: Class-based single dispatch

> For this source code

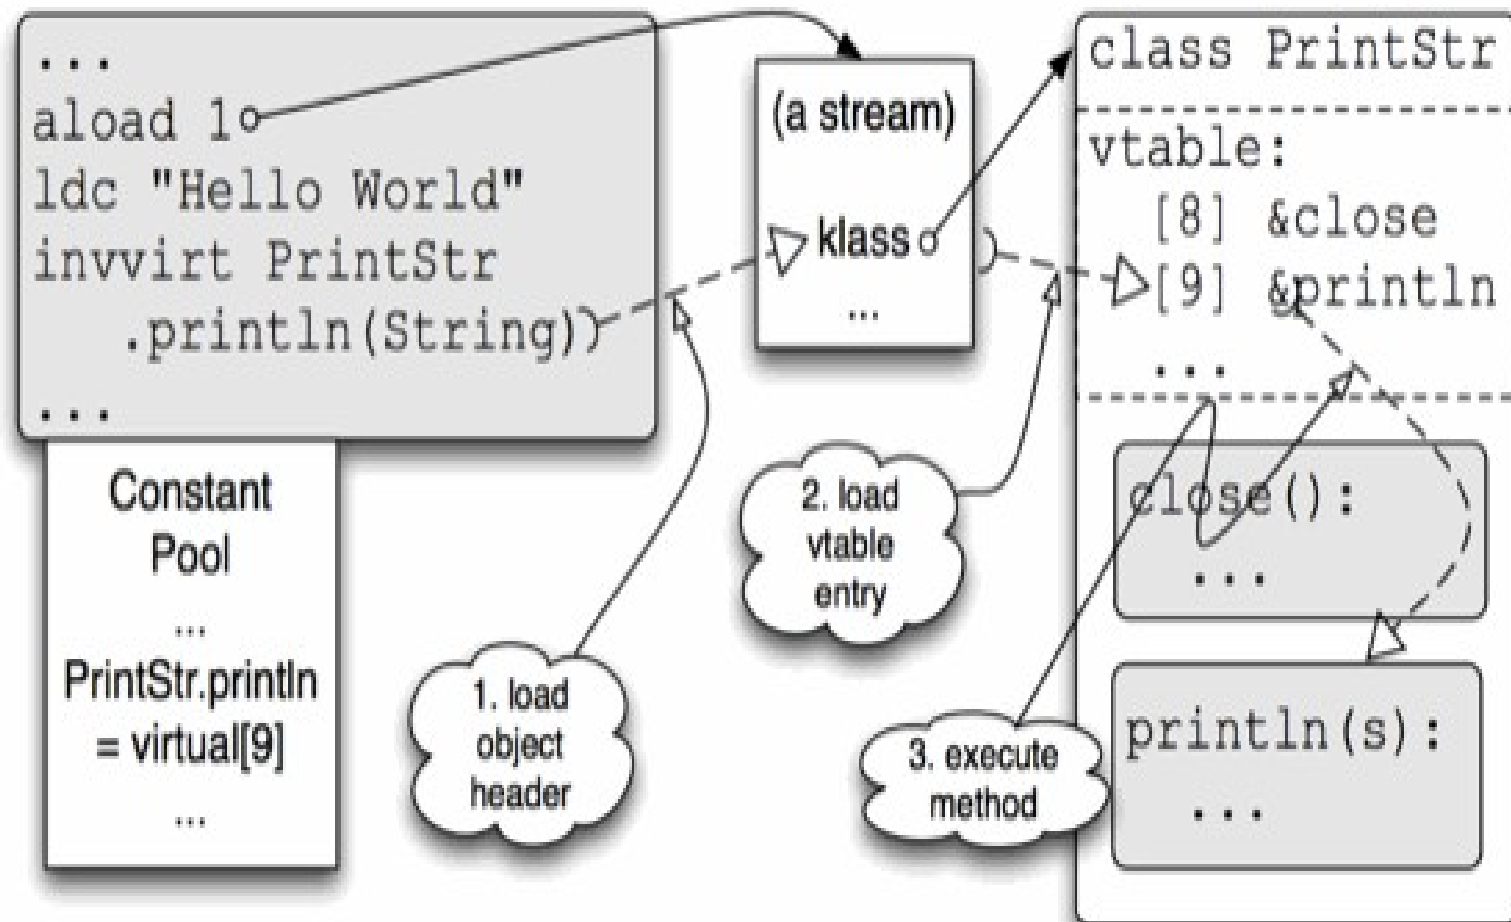
```
//PrintStream out = System.out;  
out.println("Hello World");
```

The compiled byte code looks like

```
4:  aload 1  
5:  ldc #2          //String "Hello World"  
7:  invokevirtual #4 //Method java/io/PrintStream.println:  
                        (Ljava/lang/String;)V
```

- a) Again, names in bytecode
- b) Again, linking fixed by JVM
- c) *Only* the receiver type determines method selection
- d) *Only* the receiver type can be adapted (narrowed)

How the VM selects the target method:



(Note: This implementation is typical; VMs vary.)

What more could anybody want? (1)

- > Naming — not just Java names
 - > arbitrary strings, even structured tokens (XML??)
 - > help from the VM resolving names is **optional**
 - > caller and callee do **not** need to agree on names
- > Linking — not just Java & VM rules
 - > can link a call site to any callee the runtime wants
 - > can *re-link* a call site if something changes
- > Selecting — not just static or receiver-based
 - > selection logic can look at any/all arguments
 - > (or any other conditions relevant to the language)

What more could anybody want? (2)

- > Adapting — no exact signature matching
 - > widen to Object, box from primitives
 - > checkcast to specific types, unbox to primitives
 - > collecting/spreading to/from varargs
 - > inserting or deleting extra control arguments
 - > language-specific coercions & transformations
- > (...and finally, the same fast control transfer)
- > (...with inlining in the optimizing compiler, please)

Dynamic method invocation

- > How would we compile a function like

```
function max(x, y) {  
    if (x.lessThan(y)) then y else x  
}
```

- > Specifically, how do we call `.lessThan()`?

Dynamic method invocation (how not to)

> How about:

```
0:    aload 1; aload 2
2:    invokevirtual #3    //Method Unknown.lessThan:
                                (Ljava.lang.Object;) Z
5:    if_icmpeq
```

> That doesn't work

- > No receiver type
- > No argument type
- > Return type might not even be boolean ('Z')

Dynamic method invocation (how to)

> A new option:

```
0:   aload_1; aload 2
2:   invokedynamic #3   //NameAndType lessThan:
                        (Ljava/lang/Object;Ljava/lang/Object;) Z
5:   if_icmpeq
```

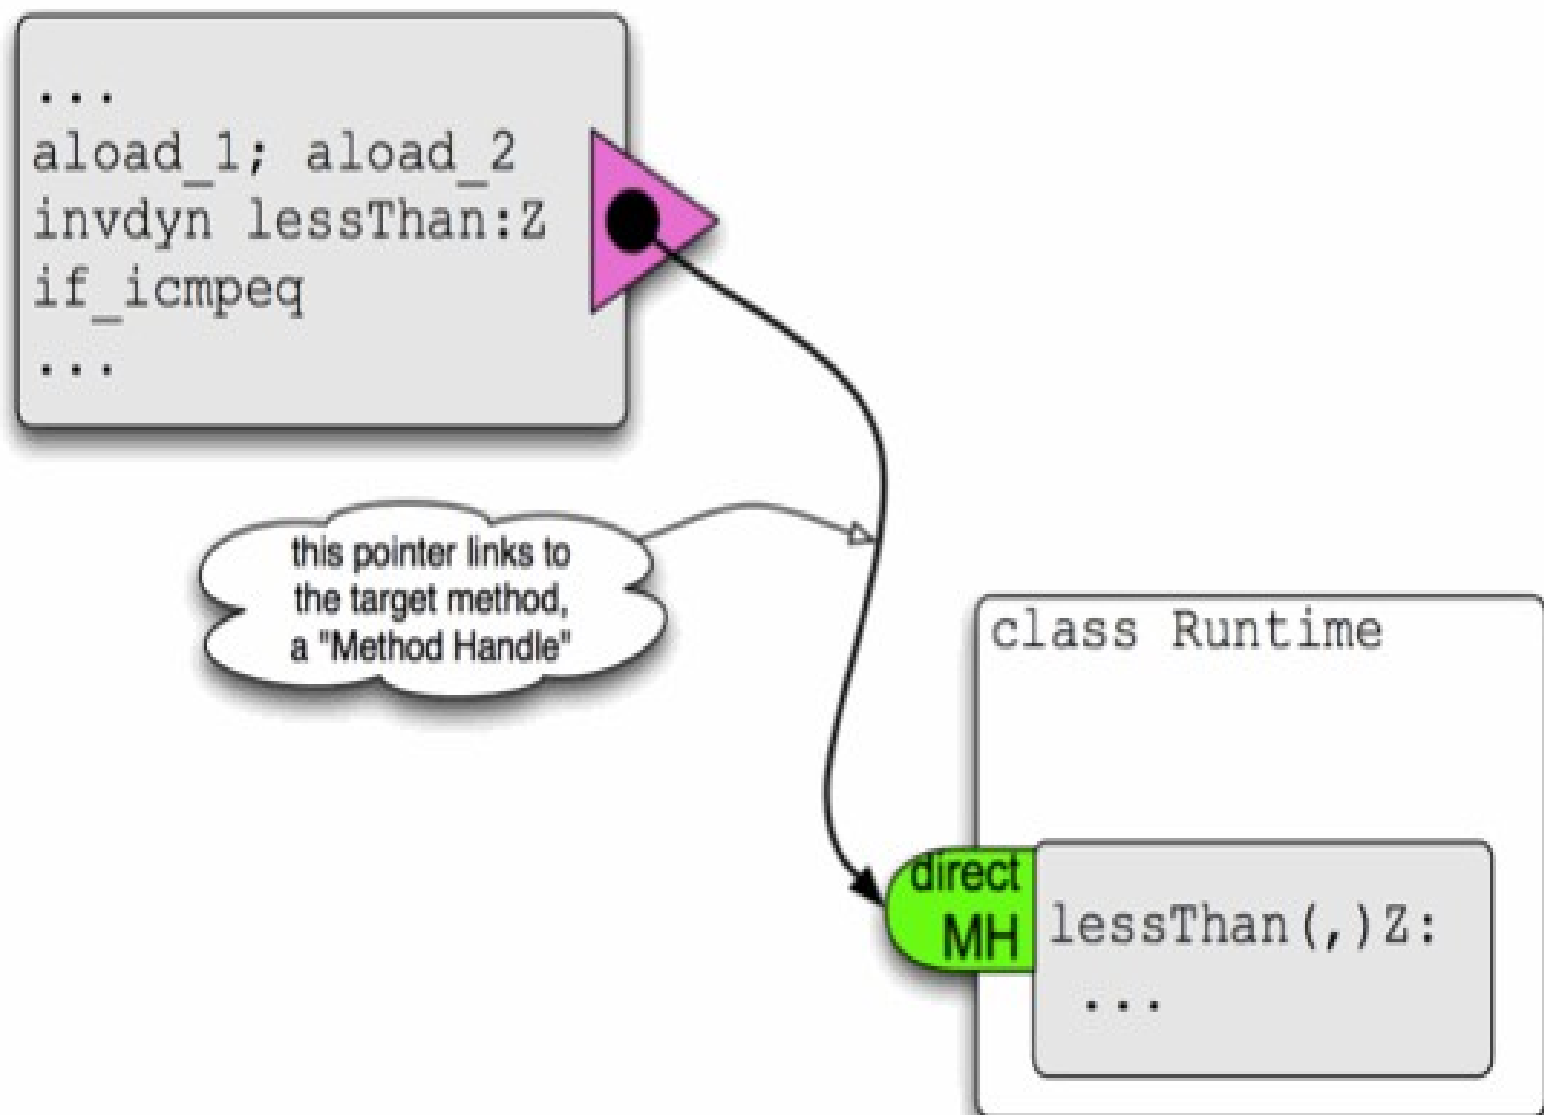
> Advantages:

- Compact representation
- Argument types are untyped Objects
- Required boolean return type is respected

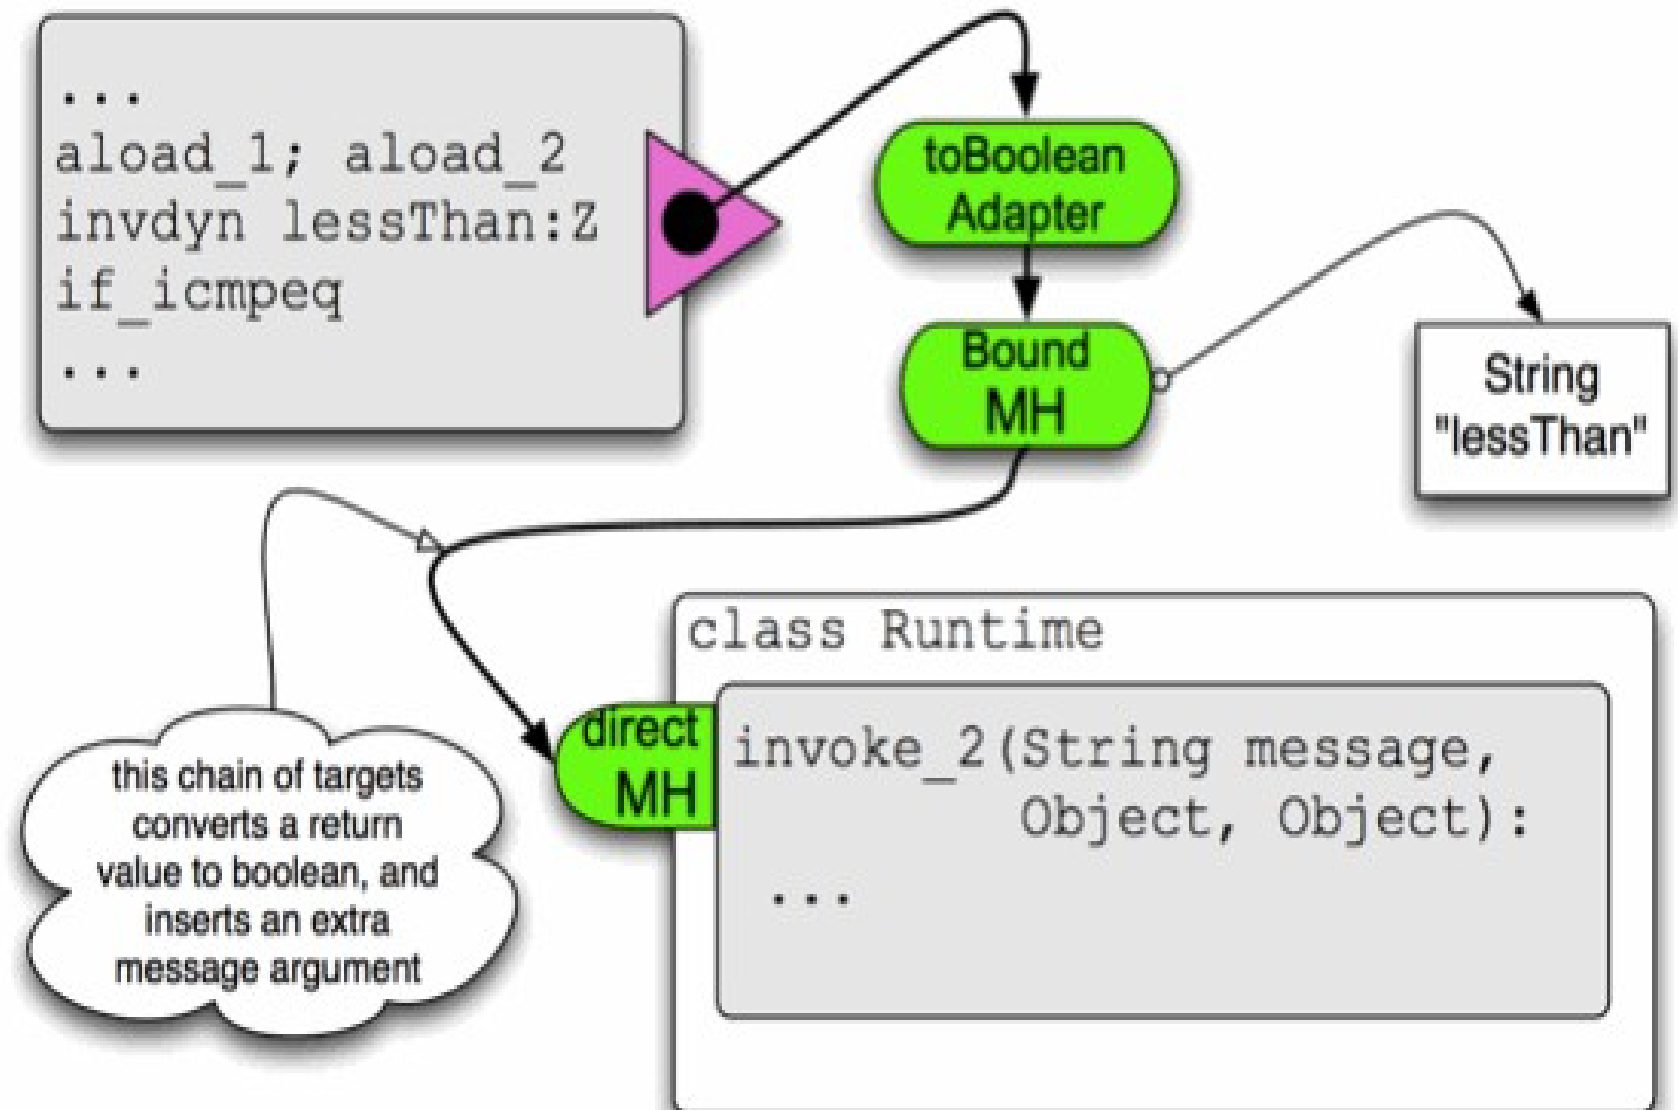
Dynamic method invocation (details)

- > But where is the dynamic language plumbing??
 - > We need something like **invoke_2** and **toBoolean!**
 - > How does the runtime know the name **lessThan**?
- > Answer: it's all *method handles* (MH).
 - > A MH can point to any accessible method
 - > (A MH can do normal receiver-based dispatch)
 - > The target of an invokedynamic is a MH

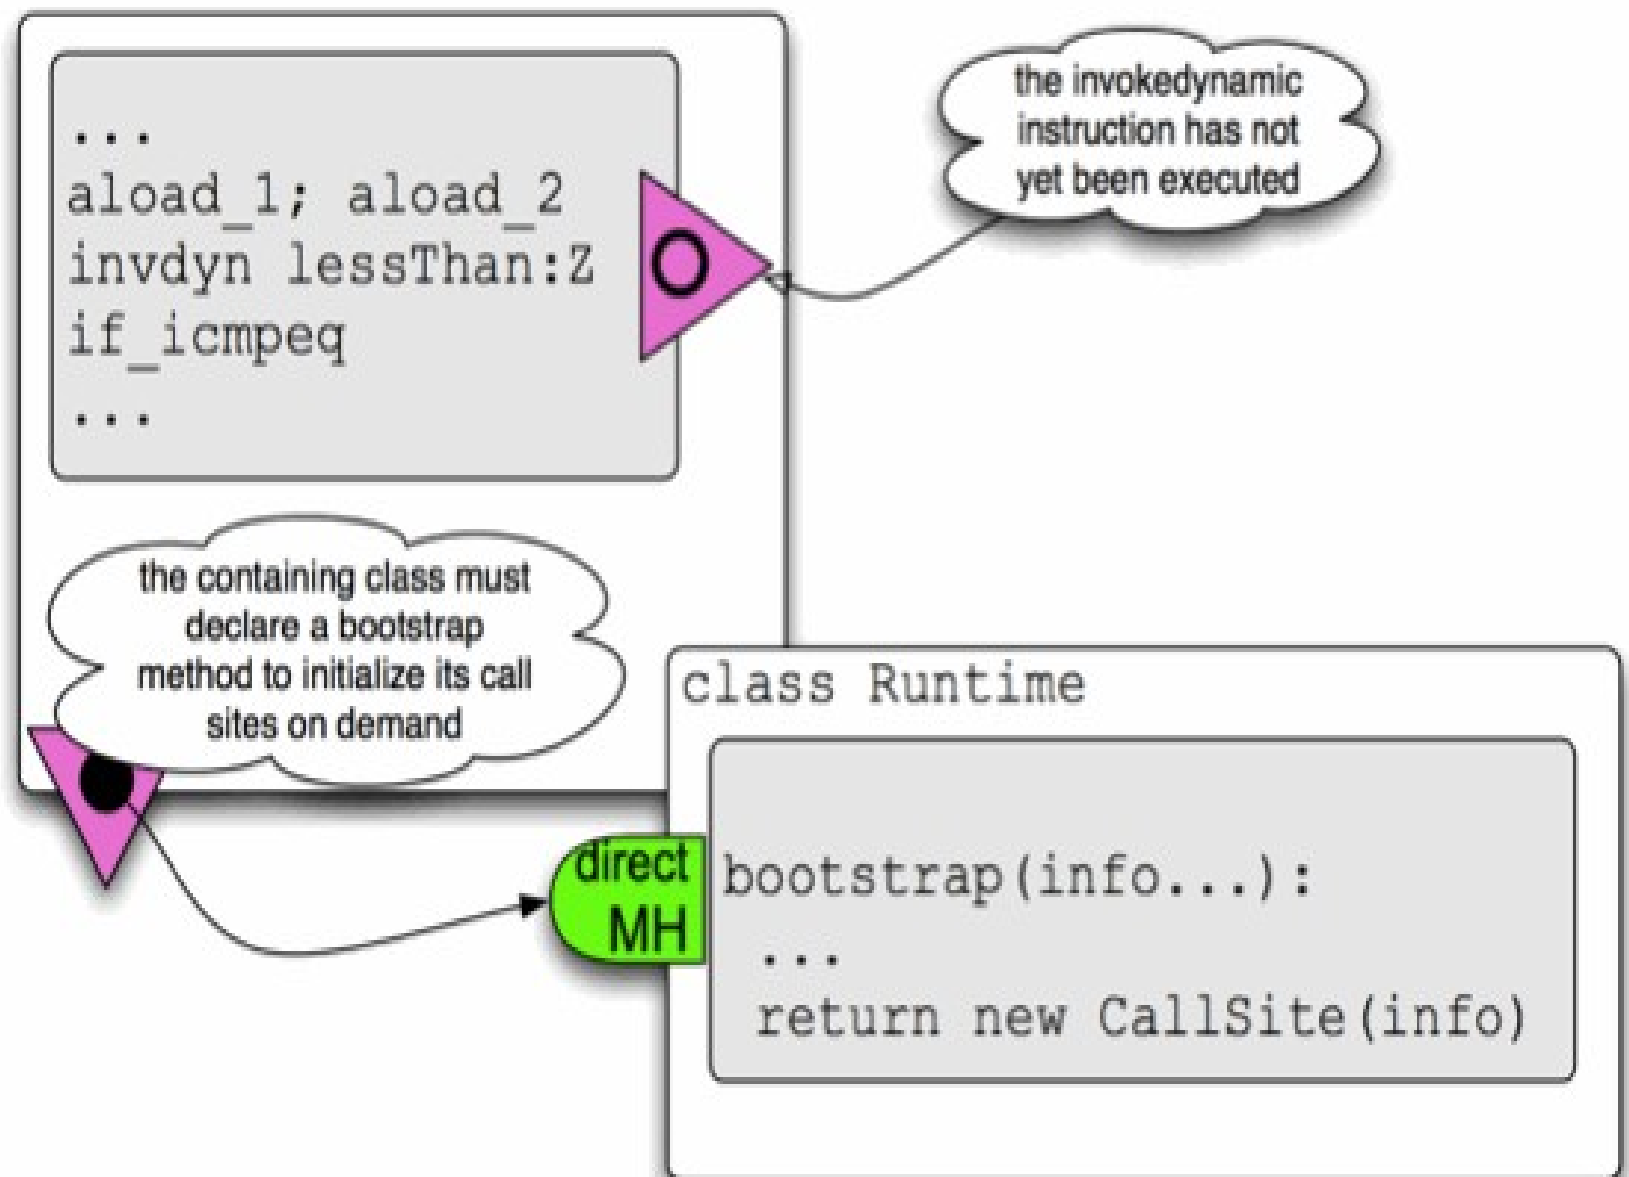
invokedynamic, as seen by the VM:



more invokedynamic plumbing: “adapters”



meta-plumbing: the bootstrap method



A Budget of Invokes

<code>invoke- static</code>	<code>invoke- special</code>	<code>invoke- virtual</code>	<code>invoke- interface</code>	<code>invoke- dynamic</code>
no receiver	receiver	receiver class	receiver interface	no receiver
no dispatch	no dispatch	single dispatch	single dispatch	adapter- based dispatch
<code>B8 nn nn</code>	<code>B7 nn nn</code>	<code>B6 nn nn</code>	<code>B9 nn nn aa 00</code>	<code>BA nn nn 00 00</code>